

Algorithms

Chapter 8

Sorting in Linear Time

Assistant Professor: Ching-Chi Lin

林清池 助理教授

chingchi.lin@gmail.com

Department of Computer Science and Engineering
National Taiwan Ocean University

Outline

- ▶ **Lower bounds for sorting**
- ▶ Counting sort
- ▶ Radix sort
- ▶ Bucket sort

Overview

- ▶ Sort n numbers in $O(n \lg n)$ time
 - ▶ Merge sort and heapsort achieve this upper bound in the worst case.
 - ▶ Quicksort achieves it on average.
 - ▶ For each of these algorithms, we can produce a sequence of n input numbers that causes the algorithm to run in $\Theta(n \lg n)$ time.
- ▶ **Comparison sorting**
 - ▶ The only operation that may be used to gain order information about a sequence is comparison of pairs of elements.
 - ▶ All sorts seen so far are comparison sorts: insertion sort, selection sort, merge sort, quicksort, heapsort.

Lower bounds for sorting

▶ **Lower bounds**


- ▶ $\Omega(n)$ to examine all the input.
- ▶ All sorts seen so far are $\Omega(n \lg n)$.
- ▶ We'll show that $\Omega(n \lg n)$ is a lower bound for comparison sorts.

▶ **Decision tree**

- ▶ Abstraction of any comparison sort.
- ▶ A full binary tree.
- ▶ Represents comparisons made by
 - ▶ a specific sorting algorithm
 - ▶ on inputs of a given size.
- ▶ Control, data movement, and all other aspects of the algorithm are ignored.

Decision tree

- ▶ For insertion sort on 3 elements:



- ▶ How many leaves on the decision tree?
 - ▶ There are $\geq n!$ leaves, because every permutation appears at least once.

Properties of decision trees_{1/3}

▶ **Lemma 1** Any binary tree of height h has $\leq 2^h$ leaves.

▶ *Proof:* By induction on h .

▶ **Basis:**

▶ $h = 0$. Tree is just one node, which is a leaf. $2^0 = 1$.

▶ **Inductive step:**

▶ Assume true for height $= h - 1$.

▶ Extend tree of height $h - 1$ by making as many new leaves as possible.

▶ Each leaf becomes parent to two new leaves.

▶ # of leaves for height $h = 2 \cdot (\text{\# of leaves for height } h - 1)$
 $= 2 \cdot 2^{h-1}$ (ind. hypothesis)
 $= 2^h$.

Properties of decision trees_{2/3}

- ▶ **Theorem 1** Any decision tree that sorts n elements has height $\Omega(n \lg n)$.

Proof:

- ▶ $\ell \geq n!$, where $\ell = \#$ of leaves.
- ▶ By lemma 1, $n! \leq \ell \leq 2^h$ or $2^h \geq n!$.
- ▶ Take logs: $h \geq \lg(n!)$.
- ▶ Use Stirling's approximation: $n! > (n/e)^n$

$$\begin{aligned} h &> \lg(n/e)^n \\ &= n \lg(n/e) \\ &= n \lg n - n \lg e \\ &= \Omega(n \lg n). \end{aligned}$$

Properties of decision trees_{3/3}

- ▶ **Corollary 1** Heapsort and merge sort are asymptotically optimal comparison sorts.

Proof:

- ▶ The $O(n \lg n)$ upper bounds on the running times for heapsort and merge sort match the $\Omega(n \lg n)$ worst-case lower bound from Theorem 1.

Outline

- ▶ Lower bounds for sorting
- ▶ **Counting sort**
- ▶ Radix sort
- ▶ Bucket sort

Counting sort

- ▶ Non-comparison sorts.
- ▶ Depends on a **key assumption**: numbers to be sorted are integers in $\{0, 1, \dots, k\}$.
- ▶ **Input**: $A[1 \dots n]$, where $A[j] \in \{0, 1, \dots, k\}$ for $j = 1, 2, \dots, n$.
Array A and values n and k are given as parameters.
- ▶ **Output**: $B[1 \dots n]$, sorted. B is assumed to be already allocated and is given as a parameter.
- ▶ **Auxiliary storage**: $C[0 \dots k]$.
- ▶ **Worst-case running time**: $\Theta(n+k)$.

The COUNTING-SORT procedure

COUNTING-SORT(A, B, k)

```
1.  for  $i \leftarrow 0$  to  $k$ 
2.      do  $C[i] \leftarrow 0$  }  $\Theta(k)$ 
3.  for  $j \leftarrow 1$  to  $length[A]$ 
4.      do  $C[A[j]] \leftarrow C[A[j]] + 1$  }  $\Theta(n)$ 
5.  /*  $C[i]$  now contains the number of elements equal to  $i$ . */
6.  for  $i \leftarrow 1$  to  $k$ 
7.      do  $C[i] \leftarrow C[i] + C[i - 1]$  }  $\Theta(k)$ 
8.  /*  $C[i]$  now contains the number of elements less than or equal to  $i$ . */
9.  for  $j \leftarrow length[A]$  downto 1
10.     do  $B[C[A[j]]] \leftarrow A[j]$ 
11.     do  $C[A[j]] \leftarrow C[A[j]] - 1$  }  $\Theta(n)$ 
```

► **The running time:** $\Theta(n+k)$.

	1	2	3	4	5	6	7	8
A	2	5	3	0	2	3	0	3



	1	2	3	4	5	6	7	8
B							3	

	0	1	2	3	4	5
C	2	0	2	3	0	1

	0	1	2	3	4	5
C	2	2	4	7	7	8

	0	1	2	3	4	5
C	2	2	4	6	7	8



	1	2	3	4	5	6	7	8
B	0	0	2	2	3	3	3	5



	1	2	3	4	5	6	7	8
B		0					3	

	1	2	3	4	5	6	7	8
B		0				3	3	

	0	1	2	3	4	5
C	1	2	4	6	7	8

	0	1	2	3	4	5
C	1	2	4	5	7	8

Properties of counting sort

- ▶ A sorting algorithm is said to be **stable** if keys with same value appear in same order in output as they did in input.
- ▶ **Counting sort is stable** because of how the last loop works.
- ▶ Counting sort will be used in radix sort.

Outline

- ▶ Lower bounds for sorting
- ▶ Counting sort
- ▶ **Radix sort**
- ▶ Bucket sort


Radix sort

- ▶ **Key idea:** Sort **least** significant digits first.

RADIX-SORT(A, d)

1. **for** $i \leftarrow 1$ **to** d
2. **do** use a stable sort to sort array A on digit i

- ▶ An example:



Correctness of radix sort

- ▶ **Proof:** By induction on number of passes (i in pseudocode).
- ▶ **Basis:**
 - ▶ $i = 1$. There is only one digit, so sorting on that digit sorts the array.
- ▶ **Inductive step:**
 - ▶ Assume digits 1, 2, ..., $i - 1$ are sorted.
 - ▶ Show that a stable sort on digit i leaves digits 1, 2, ..., i sorted:
 - ▶ If 2 digits in position i are different, ordering by position i is correct, and positions 1, ..., $i - 1$ are irrelevant.
 - ▶ If 2 digits in position i are equal, numbers are already in the right order (by inductive hypothesis). The stable sort on digit i leaves them in the right order.

Time complexity of radix sort

- ▶ Assume that we use counting sort as the intermediate sort.
- ▶ When each digit is in the range 0 to $k-1$, each pass over n d -digit number takes time $\Theta(n + k)$.
- ▶ There are d passes, so the total time for radix sort is $\Theta(d(n + k))$.
- ▶ If $k = O(n)$, time = $\Theta(dn)$.
- ▶ **Lemma 2:** Given n d -digit numbers in which each digit can take on up to k possible values, RADIXSORT correctly sorts these numbers in $\Theta(d(n + k))$ time.

Break each key into digits_{1/2}

- ▶ **Lemma 3:** Given n b -bit numbers and any positive integer $r \leq b$, RADIX-SORT correctly sorts these numbers in $\Theta((b/r)(n + 2^r))$ time.
- ▶ **Proof**
 - ▶ We view each key as having $d = \lceil b/r \rceil$ digits of r bits each.
 - ▶ Each digit is an integer in the range 0 to $2^r - 1$, so that we can use counting sort with $k = 2^r - 1$.
 - ▶ Each pass of counting sort takes time $\Theta(n+k) = \Theta(n+2^r)$.
 - ▶ A total running time of $\Theta(d(n + 2^r)) = \Theta((b/r)(n + 2^r))$.
- ▶ **For example:**
 - ▶ 32-bit words, 8-bit digits.
 - ▶ $b = 32, r = 8, d = 32/8 = 4, k = 2^8 - 1 = 255$.

Break each key into digits_{2/2}

- ▶ Recall that the running time is $\Theta((b/r)(n + 2^r))$.
- ▶ How to choose r ?
 - ▶ Balance b/r and $n + 2^r$.
- ▶ If $b < \lfloor \lg n \rfloor$, then choosing $r = b$ yields a running time of $(b/b)(n + 2^r) = \Theta(n)$.
- ▶ If $b \geq \lfloor \lg n \rfloor$, then choosing $r \approx \lg n$ gives us $\theta(\frac{b}{\lg n}(n + n)) = \theta(\frac{bn}{\lg n})$.
 - ▶ If $r > \lg n$, then 2^r term in the numerator increases faster than the r term in the denominator.
 - ▶ If $r < \lg n$, then b/r term increases, and $n + 2^r$ term remains at $\Theta(n)$.

The main reason

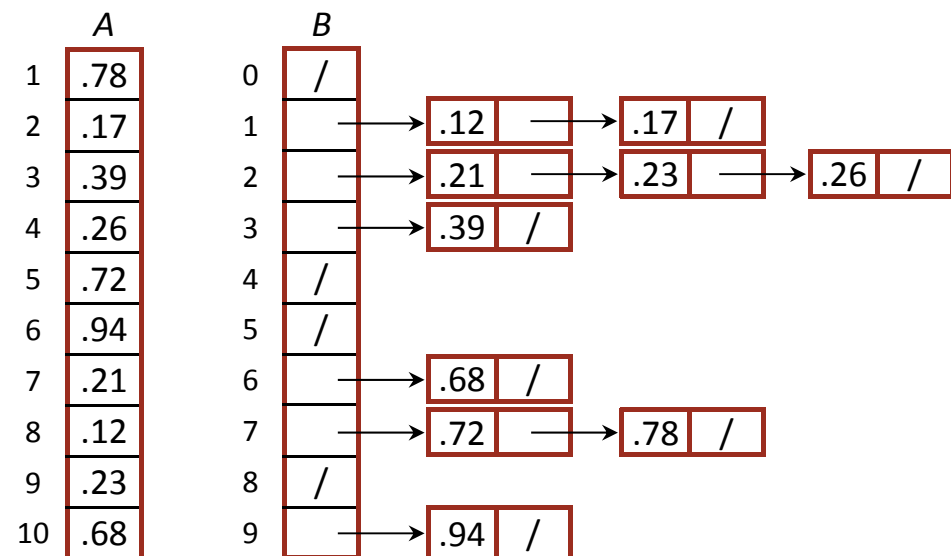
- ▶ How does radix sort violate the ground rules for a comparison sort?
 - ▶ Using counting sort allows us to gain information about keys by means other than directly comparing 2 keys.
 - ▶ Used keys as array indices.

Outline

- ▶ Lower bounds for sorting
- ▶ Counting sort
- ▶ Radix sort
- ▶ **Bucket sort**

Bucket sort

- ▶ Assumes the input is generated by a random process that distributes elements uniformly over $[0, 1)$.
- ▶ **Key idea:**
 - ▶ Divide $[0, 1)$ into n equal-sized **buckets**.
 - ▶ Distribute the n input values into the buckets.
 - ▶ Sort each bucket.
 - ▶ Then go through buckets in order, listing elements in each one.



The BUCKET SORT procedure

- ▶ **Input:** $A[1..n]$, where $0 \leq A[i] < 1$ for all i .
- ▶ **Auxiliary array:** $B[0..n-1]$ of linked lists, each list initially empty.

BUCKET-SORT(A, n)

1. **for** $i \leftarrow 1$ **to** n
2. **do** insert $A[i]$ into list $B[\lfloor n \cdot A[i] \rfloor]$
3. **for** $i \leftarrow 0$ **to** $n - 1$
4. **do** sort list $B[i]$ with insertion sort
5. concatenate lists $B[0], B[1], \dots, B[n - 1]$ together in order
6. **return** the concatenated lists

Correctness of bucket sort

- ▶ Consider $A[i], A[j]$.
Assume without loss of generality that $A[i] \leq A[j]$.
- ▶ Then $\lfloor n \cdot A[i] \rfloor \leq \lfloor n \cdot A[j] \rfloor$.
- ▶ So $A[i]$ is placed into the same bucket as $A[j]$ or into a bucket with a lower index.
- ▶ If same bucket, insertion sort fixes up.
- ▶ If earlier bucket, concatenation of lists fixes up.

Time complexity of bucket sort

- ▶ Relies on no bucket getting too many values.
- ▶ All lines of algorithm except insertion sorting take $\Theta(n)$ altogether.
- ▶ Intuitively, if each bucket gets a constant number of elements, it takes $O(1)$ time to sort each bucket $\rightarrow O(n)$ sort time for all buckets.
- ▶ We “expect” each bucket to have few elements, since the average is 1 element per bucket.

Time complexity of bucket sort

- ▶ Define a random variable: n_i = the number of elements placed in bucket $B[i]$.
- ▶ Because insertion sort runs in quadratic time, bucket sort time is $T(n) = \theta(n) + \sum_{i=0}^{n-1} O(n_i^2)$.
- ▶ Take expectations of both sides:

$$\begin{aligned} E[T(n)] &= E\left[\theta(n) + \sum_{i=0}^{n-1} O(n_i^2)\right] && \textbf{Claim that } E[n_i^2] = 2 - 1/n \text{ for } 0 \leq i \leq n-1. \\ &= \theta(n) + \sum_{i=0}^{n-1} E[O(n_i^2)] && \text{Therefore, } E[T(n)] = \theta(n) + \sum_{i=0}^{n-1} O(2 - 1/n) \\ &= \theta(n) + \sum_{i=0}^{n-1} O(E[n_i^2]) && \begin{array}{l} \text{linearity of} \\ \text{expectation} \end{array} \\ &&& \begin{array}{l} E[aX] = aE[X] \end{array} \\ &&& = \theta(n) + O(n) \\ &&& = \theta(n). \end{aligned}$$

Proof of claim

▶ **Claim:** $E[n_i^2] = 2 - 1/n$ for $0 \leq i \leq n - 1$.

▶ **Proof**

- ▶ $\Pr\{A[j] \text{ falls in bucket } i\} = p = 1/n$.
- ▶ The probability that $n_i = k$ follows the binomial distribution $b(k; n, p)$.
- ▶ So, $E[n_i] = np = 1$ and variance $\text{Var}[n_i] = np(1 - p) = 1 - 1/n$.
- ▶ For any random variable X , we have $E[n_i^2] = \text{Var}[n_i] + E^2[n_i]$

$$\begin{aligned} &= 1 - \frac{1}{n} + 1^2 \\ &= 2 - \frac{1}{n}. \end{aligned}$$

Notes

- ▶ Again, not a comparison sort. Used a function of key values to index into an array.
- ▶ This is a **probabilistic analysis**. We used probability to analyze an algorithm whose running time depends on the distribution of inputs.
- ▶ Different from a **randomized algorithm**, where we use randomization to impose a distribution.
- ▶ With bucket sort, if the input isn't drawn from a uniform distribution on $[0, 1)$, all bets are off (performance-wise, but the algorithm is still correct).